Examining the Stability of Logging Statements

Suhas Kabinna, Cor-Paul Bezemer and Ahmed E. Hassan
Software Analysis and Intelligence Lab (SAIL)
Queen's University
Kingston, Ontario
Email:{kabinna, bezemer, ahmed}@cs.queensu.ca

Weiyi Shang
Department of Computer Science and Software Engineering
Concordia University
Montreal, Quebec
Email: shang@encs.concordia.ca

Abstract—Logs are created by logging statements and they assist in understanding system behavior, monitoring choke-points and debugging. Prior research has demonstrated the importance of logging statements in operating, understanding and improving software systems. The importance of logs has lead to a new market of log management applications and tools. However, logs are often unstable, i.e., the logging statements that generate logs are often changed without the consideration of other stakeholders, causing misleading results and failure of log analysis tools. In order to proactively mitigate such issues that are caused by unstable logging statements, in this paper we empirically study the stability of logging statements in four open source applications namely: Liferay, ActiveMQ, Camel and CloudStack. We find that 20-45% of the logging statements in our studied applications change throughout their lifetime. The median number of days between the introduction of a logging statement and the first change to that statement is between 1 and 17 in our studied applications. These numbers show that in order to reduce maintenance effort, developers of log processing tools must be careful when selecting the logging statements on which they will let their tools depend.

In this paper, we make an important first step towards assisting developers in deciding whether a logging statement is likely to remain unchanged in the future. Using random forest classifiers, we examine which metrics are important for understanding whether a logging statement will change. We show that our classifiers achieve 83%-91% precision and 65%-85% recall in the four studied applications. We find that file ownership, developer experience, log density and SLOC are important metrics for deciding whether a logging statement will change in the future. Developers can use this knowledge to build more robust log processing tools, by making those tools depend on logs that are generated by logging statements which are likely to remain unchanged.

I. INTRODUCTION

Developers use logging statements to yield useful information about the state of an application during its execution. This information is collected into files (logs) and contains details which would otherwise be difficult to collect, such as the value of variables. Logs are used during various development activities such as fixing bugs [1, 2, 3], analyzing load tests [4], monitoring performance [5] and transferring knowledge [6]. Logging statements make use of logging libraries (e.g., Log4j [7]) or more archaic methods such as *print* statements. Every logging statement contains a textual part, which provides information about the context, a variable part providing context information about the event and a log

Fig. 1: An example of a logging statement

level, which shows the verbosity of the logging statement. An example of a logging statement is shown in Figure 1.

The rich knowledge in logs has lead to the development of many log processing tools such as *Splunk* [8], *Xpolog* [9], *Logstash* [10] and research tools such as Salsa [11], log-enhancer [5] and Chukwa [12] that are designed to analyze logs as well as improve logging statements. However, when logging statements are changed, the associated log processing tools may also need to be updated. For example, Figure 2 demonstrates a case in which a developer removes the elapsed time for an event. Removing information from a logging statement can affect log processing tools that rely on the removed information in order to monitor the health of the application. Prior research shows that 60% of the logging statements that generate output during system execution are changed [6]. Such changes may affect the log processing tools that heavily depend on the logs that are generated by these logging statements.

Knowing whether a logging statement is likely to change in the future is helpful to reduce the effort that is required to maintain log processing tools. If a developer of a log processing tool knows that a logging statement is likely to change, the developer can opt not to depend on the logs that are generated by this logging statement. Instead, the developer can let the log processing tool depend on output generated by logging statements that are likely to remain unchanged. Depending on logging statements that remain unchanged will reduce the maintenance effort that is required for keeping the log processing tool consistent with the ever-changing logs.

To decide whether a logging statement will change in the future, we must understand which factors play an important role during such a change. The following factors can influence whether a logging statement will change:

1) the contents of the logging statement,
2) the location of the logging statement,

Fig. 2: Modification of a logging statement

3) and the developer who added the logging statement into the source code

In this paper, we examine which of these factors can help to decide whether a logging statement will change. First, we present a preliminary study which was done to get a better understanding of the stability of logging statements in four open source applications namely ActiveMQ, Camel, Cloudstack and Liferay. In this preliminary study, we find that 20%-45% of the logging statements are changed at least once during their lifetime in the studied applications. Therefore, developers of log processing tools have to carefully select the logging statements to depend on.

Second, we explore factors that are important for explaining the stability of a logging statement using a random forest classifier. This classifier uses metrics that describe the three factors mentioned above to decide the likelihood of changing a logging statement. The most important observations in this paper are:

1) We model whether a logging statement will be changed in the future using a *random forest* classifier with 83%-91% precision and 65%-85% recall.
2) Logging statements that are added by highly experienced developers and very new developers are less likely to be changed. We find that in three of the studied applications the top three developers add more than 60% of the logging statements and 70% of the logging statements that are added by the top three developers remain untouched.
3) Logging statements added by developers who have little ownership on the file that contains the logging statements have a higher likelihood of being changed. We find that 27%-67% of all log changes, are done on logging statements added by developers who own less than 20% of the file.
4) Large files (i.e., SLOC is $2\times$-$3\times$ the median) with a low log density are more likely to have changes to their logging statements than well logged files.

The above findings can assist the maintainers of log processing tools with selecting the logging statements on which they let their tools depend. The remainder of this paper is organized as follows. Section II presents the preliminary analysis that motivates our study. Section III describes the random forest classifier and the analysis results. Section IV describes the prior research that is related to our work. Section V discusses the threats to validity. Finally, Section VI concludes the paper.

II. PRELIMINARY ANALYSIS

In this paper we study the changes that are made to logging statements in open source applications. The goal of our study is to present a classifier for deciding whether a logging statement is likely to change in the future. This classifier can assist developers of log processing tools in deciding on which logging statements they want their tool to depend. First, we perform a preliminary analysis, in which we examine how often logging statements change, to motivate our work. In this section, we present our rationale for selecting the applications that we studied and present the results of our preliminary analysis of the four studied applications.

A. Studied Applications

We selected our studied applications based on the following three criteria:

- **Usage of logging statements.** The applications must make extensive use of logging statements in their source code.
- **Project activity.** The applications must have a mature development history (i.e., more than 10,000 commits).
- **Technology used.** To simplify the implementation of our study, we opted to only select applications that are written in Java and are available through a Git repository.

To select applications that match these criteria, we first selected all Java applications from the list of Apache Foundation Git repositories[1] that have more than 10,000 commits. Next, we counted the number of logging statements in all `*.java` files in a repository using the `grep` command in Listing 1.

```
grep -icR
"\(log\.*\)\.\(|info\|trace\|debug\|error\|warn\)(" .
| grep "\.java"
```

Listing 1: Counting logging statements

This command counts the occurrences in a file of an invocation of a logging library (e.g., `log` or `_logger`) followed by the specification of a log level. We sum the occurrences in all files of an application to get the total number of logging statements shown in Table I.

We select the four applications (ActiveMQ, Camel, Cloudstack and Liferay) with the highest number of logging statements for further analysis. ActiveMQ[2] is an open source message broker and integration patterns server. Camel[3] is an open source integration platform based on enterprise integration patterns. CloudStack[4] is an open source application designed to deploy and manage large networks of virtual machines. Liferay[5] is an open source platform for building websites and web portals. Table I presents an overview of the studied applications.

[1]https://git.apache.org/
[2]http://activemq.apache.org/
[3]http://camel.apache.org/
[4]https://cloudstack.apache.org/
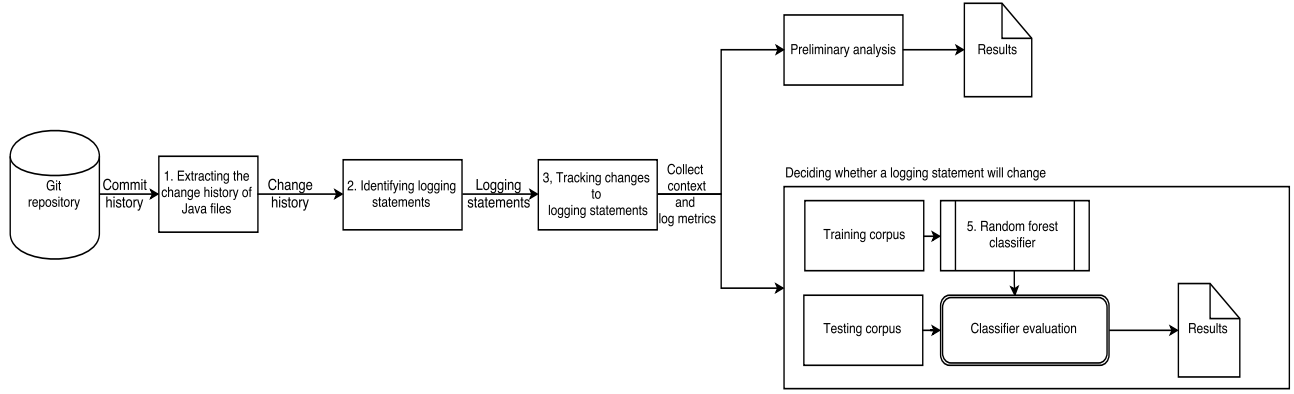[5]http://www.liferay.com/

Fig. 3: Overview of the data extraction and empirical study approach

TABLE I: An overview of the studied applications (all metrics calculated using the latest HEAD of the repository)

	ActiveMQ	Camel	CloudStack	Liferay
# of logging statements	5.1K	6.1K	9.6K	1.8K
# of commits	11K	21K	29K	143K
# of years in repository	8	8	4	4
# of Contributors	41	151	204	351
# of added lines of code	261K	505K	1.09M	3.9M
# of deleted lines of code	114K	174K	750K	2.8M
# of added logging statements	4.5K	5.1K	24K	10.4K
# of deleted logging statements	2.3K	2.4K	17K	8.1K
% of logging-related changes	1.8%	1.1%	2.3%	0.3%

B. Data Extraction Approach

The data extraction approach from the four studied applications consists of three steps, which are explained further in this section:

1) We clone the Git repository of each studied application in order to extract the change history of each file.
2) We identify the logging statements in the repository.
3) We track the changes that are made to each logging statement across commits.

We use R [13] to perform our preliminary analysis. Figure 3 shows a general overview of our approach and we detail below each of the aforementioned steps.

B.1. Extracting the change history of Java files: To examine the changes that are made to logging statements, we must first obtain a complete history of each Java file in the latest version of the main branch. We collect all the Java files in the four studied application and we use their Git repositories to obtain all the changes that are made to the files. We use Git's *follow* option to track a file even when it is renamed or relocated. We include only the changes to logging statements that are made in the main branch as other logging statements are unlikely to affect log processing tools.

B.2. Identifying logging statements: From the extracted change history of each Java file, we identify all the logging

statements. First, we manually examine the documentation of each studied application to identify the logging library used to generate the logs. We find that the studied applications use *Log4j* [14], *Slf4j*[6] and *logback*[7]. Using this information, we manually identify the common method invocations that invoke the logging library. For example, in ActiveMQ and Camel, a logging library is invoked by a method named *LOG* as shown below.

> LOG.debug(“Exception detail”, exception);

As an application can use multiple logging libraries throughout its lifetime, we use regular expressions to search for all the common log invocation patterns (i.e., *LOG, log, _logger, LOGGER, Log*). We identify every successful match of this regular expression that is followed by a log level (*info, trace, debug, error, warn*) as a logging statement.

B.3. Tracking changes to logging statements: After identifying all the logging statements, we track the changes made to these statements after their introduction. We extract the change information from the Git commits, which show a *diff* of added and removed code. To distinguish between a change in which a new logging statement is added and a change to an existing logging statement, we must track the changes made to a logging statement starting from the first application commit. Because there may be multiple changes to logging statements in a commit, we must decide to which existing logging statement a change maps.

We first collect all the logging statements in the initial commit as the initial set of logging statements. Then, we analyze the next commit to find changes to logging statements until we reach the latest commit in the repository. To distinguish between added, deleted and changed logging statements and to map the change to an existing logging statement, we use the Levenshtein ratio [15].

We use the Levenshtein ratio instead of string comparison, because the Levenshtein ratio quantifies the difference between

---

[6]http://www.slf4j.org/
[7]http://logback.qos.ch/

```
1  -         LOG.debug("Call: " +method.getName()+ " " + callTime);
2  +         LOG.debug("Call: " +method.getName()+ " took "+ callTime + "ms"); // (Statement a1)
3  +         LOG.debug("Call: " +method.setName()+ " took "+ callTime + "ms"); // (Statement a2)
```

Listing 2: Selecting the best matching logging statement

TABLE II: Distribution of changes to logging statemetns in the studied applications

Application	Unchanged logging statements	Changed logging statements
ActiveMQ	71.9 %	28.0 %
Camel	64.79 %	35.20 %
Cloudstack	55.41 %	44.58 %
Liferay	78.7 %	21.2 %

the strings on a continuous scale between 0 and 1 (the more similar the strings are, the closer the ratio approaches 1). This continuous scale is necessary to decide between multiple logging statements which can have a similar match to a change. Selecting the best matching logging statement is demonstrated by the example in Listing 2. In this example, there are two changes made to logging statements: one change and one addition.

To identify the change to logging statements, we calculate the Levenshtein ratio between each deleted and all the added logging statements and select the pair which has the highest Levenshtein ratio. This calculation is done iteratively to find all the changes within a commit. In our example, we find that the Levenshtein ratio between the deleted statement and statement *a1* is 0.86 and between the deleted statement and statement *a2* 0.76. Hence, we consider *a1* as a change. If there are no more deleted logging statements, *a2* is considered a newly added instead of a changed logging statement.

We extend the initial set of logging statements with every newly added logging statement. As we do not have change information for logging statements which are added near the end of the lifetime of the repository, we exclude these logging statements from our analysis. We find that in the studied applications, the maximum number of commits between the addition of a logging statement and its first change is 390, as shown in Figure 4. We exclude all logs added to the application 390 commits before the last commit of our analysis.

C. Results

> *Developers change 25%-45% of the logging statements in the studied applications. The median number of days between the addition of a logging statement and its first change is between 1 and 17.*

Figure **??** shows the percentage of logging statement changes in each of the studied application. We observe that A%, B%, C% and C% of the logging statements are changed in ActiveMQ, Camel, Cloudstack and Liferay, respectively, during their lifetime*[Ian says: add the number]*. This shows that logging statements change extensively throughout the


Fig. 4: Number of commits before an added logging statement is changed in the studied applications


Fig. 5: Number of days before an added logging statement is changed in the studied applications

lifetime of an application, which can affect log processing tools.

From Figure 5, we observe that 75% of the changes to logging statements are done within 145 days after the log is added. In fact, the largest median number of days between the addition of a logging statement and its first change is 17 in our studied applications. This number shows that, all too often, the changes to logging statements happen in a short time after the logging statement being added. Hence, it is important for developers of log processing tools to not depend on logging statements that are likely to change, as this dependency will require additional maintenance within a short time.

III. Deciding Whether a Logging Statement Will Change in the Future

In our preliminary analysis, we find that 25-45% of the logging statements are changed in our studied applications. These logging statement changes affect the log processing tools that depend on the logs that are generated by these statements, forcing developers spend more time on maintenance of their tools. By analyzing the metrics which can help to decide whether a logging statement will change, developers of log processing tools can reduce the effort spent on maintenance by letting their tool depend on logging statements that are likely to remain unchanged. In this section, we train a random forest classifier for deciding whether a logging statement will change in the future. We then evaluate the performance of our random forest classifier and use the classifier to understand which metrics increase the likelihood of a change to a logging statement.

A. Approach

[Ian says: Your categories of metrics should match the 3 things that you mentioned in intro! See my change below, make sure you change your text every where when you mentioned your metrics.] We use metrics that measure the content, the location and the developers of the logging statements to train the random forest classifier. *[Ian says: I only changed the sentence before this comment. Not the sentence after. Make sure you change the sentence after according to your intro.]* Context metrics measure the file context at the time of adding the logging statement. Log metrics measure details about the added logging statement. We use the Git repository to extract the context metrics and log metrics for the studied applications.

*[Ian says: Update this paragraph]*Table III defines each metric collected and the rationale behind our choice of each metric. We use context metrics as they describe the conditions in which the logging statement was added into the application. We use log metrics as they provide information about the added logging statement. These metrics benefit log processing tool developers as well as they do not need domain knowledge about the application to understand these metrics.

We build random forest classifier [16] to predict whether a logging statement will change in our studied applications. A random forest is a collection of decision trees in which the results of all trees are combined to form a generalized predictor. In our classifier, the context and log metrics*[Ian says: update]* are the explanatory variables and the dependent class variable is a boolean variable that represents whether the logging statement ever changed or not (i.e., '0' for not changed and '1' for changed).

Figure 6 provides an overview of the construction steps (C1 to C2) for building a random forest classifier and steps (A1 and A3) for analyzing the results. We use the statistical tool R to model and to analyze our data using the *RandomForest* package.

TABLE IV: Confusion Matrix

		Classified	
		Logging statement changed	Logging stsatement not changed
Actual	Logging statement changed	True positive (TP)	False negative (FN)
	Logging statement not changed	False positive (FP)	True negative (TN)

Step C1 - Removing Correlated and Redundant Metrics

Correlation analysis is necessary to remove the highly correlated metrics from our dataset [19]. Correlated metrics can severely impact the calculation of importance in the random forest classifier, as small changes to one correlated metric can affect the values of the other correlated metrics, causing large changes on dependent class variable.

We use Spearman rank correlation [20] to find correlated metrics in our data. Spearman rank correlation assesses how well two metrics can be described by a monotonic function. We use Spearman rank correlation instead of Pearson [21] because Spearman is resilient to data that is not normally distributed. We use the function *varclus* in R to perform the correlation analysis.

Figure 7 shows the hierarchically clustered Spearman ρ values in the ActiveMQ. The solid horizontal lines indicate the correlation value of the two metrics that are connected by the vertical branches that descend from it. We include one metric from the sub-hierarchies which have correlation $|\rho| > 0.7$. The dotted blue line indicates our cutoff value ($|\rho| = 0.7$). We use cutoff value of ($|\rho| = 0.7$) as used by prior research [22] to remove the correlated metrics before building our classifier.

We find that *total revision count* is highly correlated with *code churn in commit*, *log churn in commit* and *log addition*, because a file with more commits has higher chance of having a large commit with log changes, than a file with less commits. We exclude *total revision count*, *log churn in commit* and *log addition* and retain *code churn in commit* as it is a simpler metric to compute.

Correlation analysis does not indicate redundant metrics, i.e, metrics that can be explained by other explanatory metrics. The redundant metrics can interfere with the one another and the relation between the explanatory and dependent metrics is distorted. We perform redundancy analysis to remove such metrics. We use the *redun* function that is provided in the *rms* package to perform the redundancy analysis. We find after removing the correlated metrics, there are no redundant metrics.

Step C2 - Random Forest Generation

After we eliminate the correlated metrics from our datasets, we construct the random forest classifier. Random forest is a black-box ensemble classifier, which operates by constructing a multitude of decision trees on the training set and uses this to classify the testing set. From a training set of m logging statements a random sample of *n* components is selected with replacement [22] and using the *randomForest* function from the *randomForest* package in R, a random forest classifier is generated.

TABLE III: The investigated metrics in our classifier

| Dimension	Metrics	Values	Definition (d) – Rationale (r)
Context Metrics	Total revision count	Numerical	d: Total number of commits made to the file before the logging statement is added. This value is 0 for logging statements added in the initial commit but not for logging statements added overtime. r: Logging statements present in a file which is often changed, have a higher likelihood of being changed [17]. Hence, the more commits to a file, the higher the likelihood of a change to a logging statement.
	Code churn in commit	Numerical	d: The code churn of the commit in which a logging statement is added. r: The likelihood of change of logging statements that are added during large code changes, such as feature addition, can be different from that of logging statements added during bug fixes which have less code changes.
	Variables declared	Numerical	d: The number of variables which are declared before the logging statement in that function. r: When a large number of variables are declared, there is a higher chance that any of the variables will be added to or removed from a logging statement afterwards.
	SLOC	Numerical	d: The number of lines of code in the file. r: Large files have more functionality and are more prone to changes [18] and changes to logging statements [14, 17].
	File ownership	Numerical	d: Percentage of the file written by the developer who added the logging statement. r: The owner of the file is more likely to add stable logging statements than developers who have not edited the file before.
	Developer experience	Numerical	d: The number of commits the developer has made prior to this commit. r: More experienced developers may add more stable logging statements than a new developer who has less knowledge of the code.
	Log context	Categorical	d: The block in which a logging statement is added i.e., *if, if-else, try-catch, exception, throw, new function*. r: The stability of logging statements used in logical branching and assertion checks, i.e., if-else blocks, may be different from the logging statements in try-catch, exception blocks.
Log Metrics	Log addition	Boolean	d: Check if the logging statement is added to the file after creation or it was added when file was created. r: Newly added logging statements may be more likely to be changed than logging statements that exist since the creation of the file.
	Log variable count	Numerical	d: Number of logged variables. r: Over 62% of logging statement changes add new variables [17]. Hence, fewer variables in the initial logging statement might result in addition of new variables later.
	Log density	Numerical	d: Ratio of the number of logging statements to the source code lines in the file. r: Files that are well logged (i.e., with higher log density) may not need additional logging statements and are less likely to be changed.
	Log level	Categorical	d: The level (verbosity) of the added logging statement, i.e., *info,error, warn, debug, trace* and *trace*. r: Research has shown that developers spend significant amount of time in adjusting the verbosity of logging statements [17]. Hence, the verbosity level of a logging statement may affect its stability.
	Log text count	Numerical	d: Number of text phrases logged. We count all text present between a pair of quotes as one phrase. r: Over 45% of logging statements have modifications to static context [17]. Logging statements with fewer phrases might be subject to changes later to provide a better explanation.
	Log churn in commit	Numerical	d: The number of logging statements changed in the commit. r: Logging statements can be added as part of a specific change or part of a larger change.

Fig. 6: Overview of random forest classifier construction (C), analysis (A) and flow of data in random forest generation


Fig. 7: *[Ian says: I would not specify height of a figure, in order to keep the original ratio between the height and the width of figure.]*Hierarchical clustering of variables according to Spearmans ρ in ActiveMQ

Step A1 - Model Validation

After we build the random forest classifier, we evaluate the performance of our classifier using precision, recall, F-measure, AUC and Brier Score. These measures are functions of the confusion matrix as shown in Table IV and are explained below.

Precision (P) measures the correctness of our classifier in predicting which logging statement will change in the future. Precision is defined as the number of logging statements which were correctly classified as changed over all logging statements classified to have changed as explained in Equation 1.

$$P = \frac{TP}{TP + FP} \tag{1}$$

Recall (R) measures the completeness of our classifier. A classifier is said to be complete if the classifier can correctly classify all the logging statements which will get changed in our dataset. Recall is defined as the number of logging statements which were correctly classified as changed over all logging statements which actually change as explained in Equation 2.

$$R = \frac{TP}{TP + FN} \tag{2}$$

F-Measure is the harmonic mean of precision and recall, combining the inversely related measure into a single descriptive statistic as shown in Equation 3 [23].

$$F = \frac{2 \times P \times R}{P + R} \tag{3}$$

Area Under Curve (AUC) is used to measure the overall ability of the classifier to classify changed and unchanged logging statements. *[Ian says: You only say why use AUC, but what is AUC? how is AUC calculated?]* The value of AUC ranges between 0.5 (worst) for random guessing and 1 (best) where 1 means that our classifier can correctly classify every logging statement as changed or unchanged.

Brier score (BS) is a measure of the accuracy of the classifications of our classifier [24]. The Brier score explains how well the classifier performs compared to random guessing as explained in Equation 4, where P_t is the probability of whether a logging statement will change and O_t is the boolean that shows whether the statement actually changed. A perfect classifier will have a Brier score of 0, a perfect misfit classifier will have a Brier score of 1 (predicts probability of log change when log is not changed) and for random guessing the Brier score is 0.25. As a result, the lower the Brier score value, the better our random forest classifier performs.

$$BS = (P_t - O_t)^2 \tag{4}$$

Optimism: The performance measures described previously may overestimate the performance of the classifier due to overfitting. To account for the overfitting in our classifier, we use the *optimism* measure, as used by prior research [22]. The *optimism* of the performance measures are calculated as follows:

1) From the original dataset with *m* records, we select a bootstrap sample with *n* records with replacement.
2) Build random forest as described in (C2) using the bootstrap sample.
3) Apply the classifier built from the bootstrap sample on both the bootstrap and original data sample, calculating precision, recall, F-measure and Brier score for both data samples.

4) Calculate the *optimism* by subtracting the performance measures of the bootstrap sample from the original sample.

The above process is repeated 1,000 times and the average (mean) *optimism* is calculated. Finally, we calculate *optimism-reduced* performance measures for precision, recall, F-measure, AUC and Brier score by subtracting the averaged optimism of each measure, from their corresponding original measure. The smaller the optimism values, the less the original classifier overfits the data.

Step A2 - Identifying Important Metrics

To find the importance of each metric in a random forest classifier, we use a permutation test. In this test, the classifier built using the bootstrap data (i.e., two thirds of the original data) is applied to the test data (i.e., remaining one third of the original data). Then, the values of the X_i th metric of which we want to find importance for, are randomly permuted in the test dataset and the precision of the classifier is recomputed. The decrease in precision as a result of this permutation is averaged over all trees, and is used as a measure of the importance of metric X_ith in the random forest.

We use the *importance* function defined in *RandomForest* package of R, to calculate the importance of each metric. We call the *importance* function every time during the bootstrapping process to obtain 1,000 importance scores for each metric in our dataset.

As we obtain 1,000 data sets for each metric from the bootstrapping process, we use the **Scott-Knott Effest size clustering** (SK-ESD) to group the metric based on their effect size [25]. Such an approach groups metrics based on their importance in predicting the likelihood of logging statement changes. The SK-ESD algorithm uses effect size calculated using *cohen's* delta [26], to merge any two statistically different groups. This assures the means of metrics within a group not to be statistically significantly different*[Ian says: This is wrong. What you said is the original Scott-knott. The SK-ESD makes sure the the effect sizes is negligible (based on the threshold that you wrote in the next sentence), or not statistically significant.]*. We use the *SK.ESD* function in the *ScottKnottESD* package of R and set the effect size threshold parameter to negligible, (i.e., < 0.2) to cluster the two metrics into the same groups.

Step A3 - Plotting the Important Metrics

To understand the effect of each metric in our random forest classifier, it is necessary to plot the predicted probabilities of a change to a logging statement against the metrics. By plotting the predicted probabilities of a change to a logging statement, we obtain a clearer picture of how the random forest classifier uses the important metrics to classify the data.

Using the *randomForest* package in R, we build a classifier as explained in C2, and we use the *predict* function in *R*, to calculate the probabilities of a change to a logging statement. We plot each predicted probability against the value of the

TABLE VI: Contribution of top 3 developers

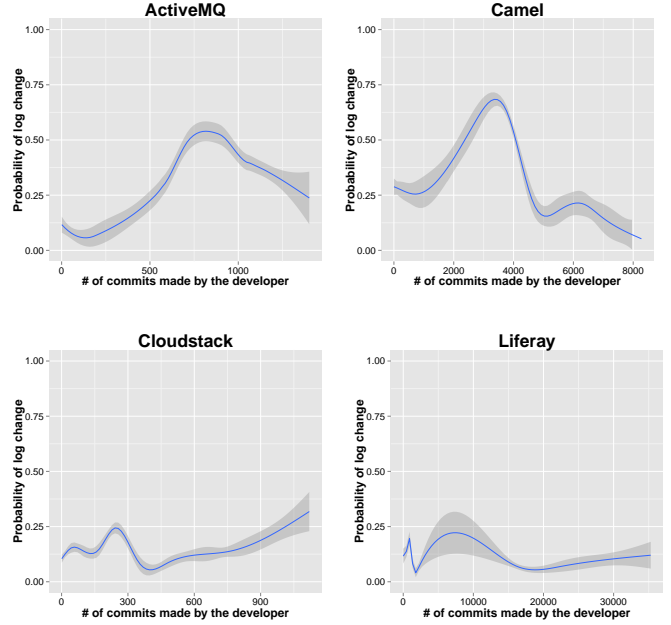|	Total logs	Changed logging statements	Total # of contributors
ActiveMQ	956 (50.4%)	301 (31.4%)	41
Camel	3060 (63.1%)	1460 (47.7%)	151
Cloudstack	5982 (35.7%)	2276 (38.0%)	204
Liferay	3382 (86.7%)	609 (18.0%)	351
Average	**3345 (59%)**	**1161 (33.75%)**	**747**

Fig. 9: Comparing the probability of a change to a logging statement against the experience of the developer who adds that logging statement

metric, to understand how changes in the metric values affect the probability of a change to a logging statement.

B. Results

> *The random forest classifier achieves 0.89-0.91 precision, 0.71-0.83 recall and outperforms random guessing for our studied applications.*

Figure 8 shows the optimism-reduced values of *precision*, *recall*, *F-measure AUC* and *Brier score* for each studied application. The classifier achieves an AUC of 0.94-0.95 and Brier scores between 0.04 and 0.07 across all studied applications. Using the equation described above, we can calculate that a Brier score of 0.07 means that our model can decide with 73% probability whether a logging statement will change. The Brier score is 0.25 for random guessing (i.e., predicted value is 50%).

Fig. 8: The optimism reduced performance measures of the four applications

TABLE V: The most important metrics, divided into homogeneous groups by Scott-Knott clustering

	ActiveMQ				**Camel**	
Rank	**Factors**	**Importance**	**Rank**	**Factors**	**Importance**
1	Developer experience	0.246	1	Developer experience	0.272
2	Ownership of file	0.175	2	Ownership of file	0.151
3	Log density	0.163	3	Log level	0.138
4	Log variable count	0.101	4	SLOC	0.112
5	Log level	0.063	5	Log addition	0.090
6	Variable declared	0.048		Log density	0.088
7	Log context	0.069	6	Log variable count	0.063
8	Log text length	0.022	7	Log context	0.052
			8	Variable declared	0.051
	CloudStack			**Liferay**	
Rank	**Factors**	**Importance**	**Rank**	**Factors**	**Importance**
1	Log density	0.224	1	Log density	0.192
2	Ownership of file	0.215		Developer experience	0.195
3	SLOC	0.192	2	Ownership of file	0.190
4	Developer experience	0.182		SLOC	0.188
5	Log text length	0.120	3	Log variable count	0.162
6	Log variable count	0.115	4	Log level	0.148
7	Log level	0.102	5	Log context	0.091
8	Variable declared	0.092	6	Variable declared	0.080
9	Log context	0.061	7	Log text length	0.071

B2. Important Metrics for Logging Statement Stability

> *In three out of four studied applications, the top three developers were responsible for adding over 50% of the logging statements. Upto 70% of these statements never change.*

Table V shows the metrics that are important for deciding whether a logging statement will change in the future. From Table V, we see that developer experience is in the top four metrics for all studied applications to help explain the likelihood of a logging statement being changed in all the studied applications. Figure 9 shows the probabilities of a logging statement being changed as developer experience increases. In all the studied applications logging statements that are added by new developers have a lower probability of being

changed, when compared to those added by more experienced developers. We account this phenomenon to the fact that inexperienced developers add a very small fraction of the logs (12% in Cloudstack and 1-3% in the other applications).

We also observe that as developers get more experience the probability of a change to a logging statement decreases in ActiveMQ, Camel and Liferay. This downward trend may be explained by the fact that in ActiveMQ, Camel and Liferay, the top three developers are responsible for adding more than 50% of the logging statements as seen in Table VI. In addition, we find that upto 70% of the logging statements added by these top developers never change. These findings suggest that developers of log processing tools should let their tools depend on logging statements added by developers with a high level of experience.

Fig. 10: Comparing the probability of a change to a logging statement against ownership of the file in which the logging statement is added

> *Logging statements that are added into a file by developers who own more than 75% of that file are unlikely to be changed in the future.*

From Table V, we see that ownership of the file is in the top two metrics to help explain the likelihood of a change to a logging statement in all the studied applications. From Figure 10, we observe in all the applications that logging statements introduced by developers who own more than 75% of the file are less likely to be changed. We also observe that developers who own less than 20% of the file are responsible for 27%-67% of the changes to logging statements in the studied applications, which is seen as upward trend from 0 to 0.20 in Figure 10. These results suggest that developers of log processing tools should be more cautious when using a logging statement written by a developer who has contributed less than 20% of the file.

> *Logging statements in files with a low log density are more likely to change than logging statements in files with a high log density.*

From Table V, we observe that log density has the highest importance in Liferay and Cloudstack. We find that in these two applications, changed logging statements are in files that have a lower log density than the files containing unchanged logging statements. When we measure the median file sizes, we find that logging statements which change more are present in files with significantly higher SLOC (2x-3x higher). This suggests that large files which are not well logged are more likely to have unstable logging statements, than well logged files.

> *File ownership, SLOC, developer experience and log density are important metrics for deciding whether a logging statement will change in the future.*

IV. Related Work

In this section, we present prior empirical research done on logs and tools developed to assist in logging.

A. Log Tools

Tan et al. [11] propose a tool named SALSA, which constructs state-machines from logs. The state-machines are further used to detect anomalies in distributed computing platforms. Yuan et al. [5] show that logs need to be improved by providing additional . Their tool named *Log Enhancer* can automatically provide additional control and data flow parameters into the logs thereby improving the logs. *Log Enhancer* can improve the quality of log added and mitigate the need for changes later. However, it does not provide any insight into why some logging statements are more likely to be changed. Our paper, tries to classify which logging statements have higher likelihood of being changed later and avoid such logging statements from the analysis.

B. Empirical Studies on Logging statements

Prior research performs an empirical study on the characteristics of logging statements. Yuan et al. [17] study the logging characteristics in four open source systems. They find that over 33% of all changes to logging statements are after-thoughts and that logging statements are changed 1.8 times more often than regular code. Fu et al. [27] performed an empirical study on where developers put logging statements. They find that logging statements are used for assertion checks, return value checks, exceptions, logic-branching and observing key points. The results of the analysis were evaluated by professionals from the industry and an F-measure of over 95% was achieved.

Research also shows that logs are a source of information about the execution of large software systems for developers and end users. Shang et al. performed an empirical study on the evolution of both static logs and logs outputted during run time [14, 28]. They find that logging statements are co-evolving with software systems. However, logging statements are often modified by developers without considering the needs of operators which even affects the log processing tools which run on top of the logs produced by these statements. They highlight the fact that there is a gap between operators and developers of software systems, especially in the leverage of logs [29]. Furthermore, Shang et al. [30] find that understanding logs is challenging. They examine user mailing lists from three large open-source projects and find that users of these systems have various issues in understanding logs outputted by the system.

The existing empirical studies on logging statements show that 1) logs are leveraged by developers for different purposes and 2) logging statements are changed extensively by developers without consideration of other stakeholders, which affect practitioners and end users. These findings highlight the need

for better understanding of the stability of logging statements in the applications.

V. THREATS TO VALIDITY

In this section, we present the threats to the validity to our findings.

External Validity

Our empirical study is performed on Liferay, ActiveMQ, Camel and CloudStack. Though these studied applications have years of history and large user bases, these applications are all Java-based. Other languages may not use logging statements as extensively. Our applications are all open source and we do not verify the results on any commercial platform applications. More studies on other domains and commercial platforms, with other programming languages are needed to see whether our findings can be generalized.

Construct Validity

Our heuristics to extract logging source code may not be able to extract every logging statement in the source code. Even though the studied applications leverage logging libraries to generate logs at run-time, there may still exist user-defined logs. By manually examining the source code, we believe that we extract most of the logging statements. Evaluation on the coverage of our extracted logging statements can address this threat.

In our model, we use the first change after the introduction of a logging statement only. While the first change is sufficient for deciding whether a logging statement will change, we need more information to decide how likely it is that a logging statement will change. In future work, we will extend our model to give more specific details about a future change.

Internal Validity

Our study is based on the data obtained from Git for all the studied applications. The quality of the data contained in the repositories can impact the internal validity of our study. For example, merging commits or rewriting the history of the repository (i.e., by *rebasing* the history) may affect our results.

Our analysis of the relationship between metrics that are important factors in predicting the stability of logging statements cannot claim causal effects, as we are investigating correlation and not causation. The important factors from our random forest models only indicate that there exists a relationship which should be studied in depth in future studies.

VI. CONCLUSION

Logging statements are snippets of code, added by developers to yield valuable information about the execution of an application. Logging statements generate their output in logs, which are used by a plethora of log processing tools to assist in software testing, performance monitoring and system state comprehension. These log processing tools are completely dependent on the logs and hence are affected when logging statements are changed.

In order to reduce the effort required for the maintenance of log processing tools, we examine changes made to logging statements in four open source applications. The goal of our work is to help developers of log processing tools decide whether a logging statement is likely to change in the future. We consider our work an important first step towards helping developers to build more robust log processing tools, as knowing whether a log will change in the future allows developers to let their log processing tools rely on logs generated by logging statements that are likely to remain unchanged. The highlights of our work are:

- We find that 20%-40% of logs are changed at-least once.
- Our random forest classifier for predicting whether a log will change achieves a precision of 83%-91% and recall of 65%-85%.
- We find that log density, SLOC, developer experience, file ownership are important predictors of log stability in the studied applications.

Our findings highlight that we can correctly classify the likelihood of a log change, when log is added into the application. The important metrics from the classifier help in determining the likelihood of a log change, and developers can use this knowledge to be more selective when importing logging statements into their processing tools.

REFERENCES

[1] W. Xu, L. Huang, A. Fox, D. Patterson, and M. I. Jordan, "Detecting large-scale system problems by mining console logs," in *ACM SIGOPS '09: Proceedings of the 22nd Symposium on Operating Systems Principles*, New York, NY, USA, pp. 117–132.

[2] J.-G. Lou, Q. Fu, S. Yang, Y. Xu, and J. Li, "Mining invariants from console logs for system problem detection." in *USENIX Annual Technical Conference*, 2010.

[3] Q. F. J. L. Y. Wang and J. Li., "Execution anomaly detection in distributed systems through unstructured log analysis." in *ICDM 09: In Proceedings of 9th IEEE International Conference on Data Mining*.

[4] H. Malik, H. Hemmati, and A. Hassan, "Automatic detection of performance deviations in the load testing of large scale systems," in *Software Engineering (ICSE), 2013 35th International Conference on*, May 2013, pp. 1012–1021.

[5] D. Yuan, J. Zheng, S. Park, Y. Zhou, and S. Savage, "Improving software diagnosability via log enhancement," *ACM Transactions on Computer Systems*, vol. 30, pp. 4:1–4:28, Feb. 2012.

[6] W. Shang, Z. M. Jiang, B. Adams, A. E. Hassan, M. W. Godfrey, M. Nasser, and P. Flora, "An exploratory study of the evolution of communicated information about the execution of large software systems," *Journal of Software: Evolution and Process*, vol. 26, no. 1, pp. 3–26, 2014.

[7] Log4j. [Online]. Available: http://logging.apache.org/log4j/2.x/

[8] D. Carasso, "Exploring splunk," *published by CITO Research, New York, USA, ISBN*, pp. 978–0, 2012.

[9] Xpolog. [Online]. Available: http://www.xpolog.com/.

[10] X. Xu, I. Weber, L. Bass, L. Zhu, H. Wada, and F. Teng, "Detecting cloud provisioning errors using an annotated process model," in *Proceedings of the 8th Workshop on Middleware for Next Generation Internet Computing.* ACM, 2013, p. 5.

[11] J. Tan, X. Pan, S. Kavulya, R. Gandhi, and P. Narasimhan, "Salsa: Analyzing logs as state machines," in *WASL'08: Proceedings of the 1st USENIX Conference on Analysis of System Logs.* USENIX Association, 2008, pp. 6–6.

[12] J. Boulon, A. Konwinski, R. Qi, A. Rabkin, E. Yang, and M. Yang, "Chukwa, a large-scale monitoring system," in *Proceedings of CCA*, vol. 8, 2008, pp. 1–5.

[13] R. Ihaka and R. Gentleman, "R: a language for data analysis and graphics," *Journal of computational and graphical statistics*, vol. 5, no. 3, pp. 299–314, 1996.

[14] W. Shang, M. Nagappan, and A. E. Hassan, "Studying the relationship between logging characteristics and the code quality of platform software," *Empirical Software Engineering*, vol. 20, no. 1, pp. 1–27, 2015.

[15] M. Mednis and M. K. Aurich, "Application of string similarity ratio and edit distance in automatic metabolite reconciliation comparing reconstructions and models," *Biosystems and Information technology*, vol. 1, no. 1, pp. 14–18, 2012.

[16] J. Albert and E. Aliu, "Implementation of the random forest method for the imaging atmospheric cherenkov telescope {MAGIC}," *Nuclear Instruments and Methods in Physics Research Section A: Accelerators, Spectrometers, Detectors and Associated Equipment*, vol. 588, no. 3, pp. 424 – 432, 2008. [Online]. Available: http://www.sciencedirect.com/science/article/pii/S0168900207024059

[17] D. Yuan, S. Park, and Y. Zhou, "Characterizing logging practices in open-source software," in *ICSE '12: Proceedings of the 34th International Conference on Software Engineering.* IEEE Press, 2012, pp. 102–112.

[18] D. Zhang, K. El Emam, H. Liu *et al.*, "An investigation into the functional form of the size-defect relationship for software modules," *Software Engineering, IEEE Transactions on*, vol. 35, no. 2, pp. 293–304, 2009.

[19] J. Cohen, P. Cohen, S. G. West, and L. S. Aiken, *Applied multiple regression/correlation analysis for the behavioral sciences.* Routledge, 2013.

[20] J. H. Zar, *Spearman Rank Correlation.* John Wiley & Sons, Ltd, 2005. [Online]. Available: http://dx.doi.org/10.1002/0470011815.b2a15150

[21] R. J. Serfling, *Approximation theorems of mathematical statistics.* John Wiley & Sons, 2009, vol. 162.

[22] S. McIntosh, Y. Kamei, B. Adams, and A. E. Hassan, "An empirical study of the impact of modern code review practices on software quality," *Empirical Software Engineering*, p. To appear, 2015.

[23] G. Hripcsak and A. S. Rothschild, "Agreement, the f-measure, and reliability in information retrieval," *Journal of the American Medical Informatics Association*, vol. 12, no. 3, pp. 296–298, 2005.

[24] D. S. Wilks, *Statistical methods in the atmospheric sciences.* Academic press, 2011, vol. 100.

[25] C. Tantithamthavorn, S. McIntosh, A. E. Hassan, and K. Matsumoto, "An Empirical Comparison of Model Validation Techniques for Defect Prediction Model," http://sailhome.cs.queensu.ca/replication/kla/model-validation.pdf, 2015, under Review at Transactions on Software Engineering (TSE).

[26] V. B. Kampenes, T. Dybå, J. E. Hannay, and D. I. Sjøberg, "A systematic review of effect size in software engineering experiments," *Information and Software Technology*, vol. 49, no. 11, pp. 1073–1086, 2007.

[27] Q. Fu, J. Zhu, W. Hu, J.-G. L. R. Ding, Q. Lin, D. Zhang, and T. Xie, "Where do developers log? an empirical study on logging practices in industry," in *ICSE Companion '14: Proceedings of the 36th International Conference on Software Engineering,*, pp. Pages 24–33.

[28] W. Shang, Z. M. Jiang, B. Adams, A. E. Hassan, M. W. Godfrey, M. Nasser, and P. Flora, "An exploratory study of the evolution of communicated information about the execution of large software systems," *Journal of Software: Evolution and Process*, vol. 26, no. 1, pp. 3–26, 2014.

[29] W. Shang, "Bridging the divide between software developers and operators using logs," in *ICSE '12 :Proceedings of the 34th International Conference on Software Engineering.*

[30] W. Shang, M. Nagappan, A. E. Hassan, and Z. M. Jiang, "Understanding log lines using development knowledge," in *ICSME '14: Proceedings of the International Conference on Software Maintenance and Evolution,.* IEEE, 2014, pp. 21–30.