

Understanding the Stability of Logs in Software

SuhasKabinna · Weiyi Shang · Cor-Paul
Bezemer · Ahmed E. Hassan

Received: date / Accepted: date

Abstract Logs are system generated outputs, inserted by developers into the code. They help in understanding the system behavior, monitoring choke-points and in debugging. Recent research has demonstrated the importance of logs in operating, understanding and improving software systems. This has lead to the development of semi-automated log management applications and log processing tools. However, logs may change in time due to debugging, improvement or addition of new features. When logs are changed, the changes have to be communicated to operators and administrators. In addition, the log processing and management applications have to be updated. In this paper we study the stability of logs on four large software systems namely: Liferay, ActiveMQ, Camel and CloudStack. We conduct a manual analysis and find that 20-80% of the logs are changed in these software systems. We find four different types of log changes namely 1) text modification, 2) variable modification, 3) log level change and 4) log relocation. After characterizing the different types of log changes we build models to predict if a log added to a file will change in the future. We use code churn, log churn and developer metrics to build a random forest classifier. Our classifier can predict which logs will change in the future with 89-93 % precision, and 76-92% recall. We find that code, log and developer metrics are strong predictors of log stability in our models and can help identify which logs are more likely to change in the future. On one hand, this can help developers of log processing tools to develop more robust applications. On the other hand, system administrators can know before hand which logs are more likely to cause issues in the log processing tools and this can reduce their maintenance costs.

Suhas. Kabinna, Weiyi. Shang, Cor-Paul. Bezemer and Ahmed E.Hassan
Software Analysis and Intelligence Lab (SAIL)
Queen's University
Kingston, Ontario
E-mail: {kabinna,swy,bezemer,ahmed}@cs.queensu.ca

1 Introduction

Logs are leveraged by developers to record useful information during the execution of a system. Loggings are recorded during various developmental activities such as bug fixing [1–3], performing improvement tasks [4], monitoring performance [5] and for knowledge transfer [6]. Logging can be done through use of log libraries or archaic methods such as *print* statements. Every log contains a textual part, which provides information about context, a variable part providing information about the event and a log level, which shows the verbosity of the logs. An example of a log is shown below where *info* is the logging level, *Testing Connection to Host Id* is the event and the value of variable *host* is the variable part.

```
LOG.info( "Testing Connection to Host Id:" + host);
```

The rich and unified format of logs has lead to the development of many log processing tools such as *Splunk*, *Xpolog*, *Logstash* and in-house tools. These log processing tools are used to generate information for capacity planning of large-scale systems [7,8], to monitor system health [9] or to detect abnormal system behavior [10]. These applications rely heavily on the log messages themselves and require continuous maintenance when the format or content of logs are changed.

Research shows that only 40% of the logs at execution level, stays the same across releases and the impact of 15-80% of the changes can be minimized through robust analysis [6]. These log changes, can affect the log processing tools which heavily depend on them and maintenance cost will be high. In this paper, we track the changes made to logs across multiple releases in the four studied systems. In order to get a better understanding of the log changes we focus our research on the following RQ's.

RQ1: How much do logs change over time and why do the changes occur?

Based on our quantitative analysis of the studied systems we identify three categories of change frequency in logs. If a log is changed more than four times we categorize it as '*Frequently Changed*', if it has only three changes or less it is categorized as '*Changed*' and if there are no changes made we categorize it as '*Never Changed*'. We find that 20-80% of all logs are changed at least once throughout the lifespan of our studied systems. We find developers change logs for four main reasons namely: 'change of log level', 'text modification', 'variable modification' and 'log relocation'.

RQ2: Can code, log and developer metrics help in explaining the stability of logs?

We find that code, log and developer related metrics help in building models which can predict which logs are more likely to change in the future. We use the data from three dimensions namely, code, log and developers. Our *random forest* achieved an accuracy of 89% to 93% in all studied systems with recall of 76% to 92%, when predicting which logs have higher likelihood of getting changed. We also identify significant metrics from each dimension, that

affect the stability of logs. We find several metrics (e.g., developer experience, source lines of code, # of comments, log text length) are strong predictors in predicting if a log will change in the future.

These results show that code, log and developer related metrics can help in identifying unstable logs in our studied systems. This can help in reducing the effort needed in the maintenance of log processing applications, as system maintainers can flag the logs that have the potential of being changed in subsequent releases and track them.

The rest of this paper is organized as follows. Section 2 presents the methodology for gathering and extracting data for our study. Section 3 presents the case studies and the results to answer the two research questions. Section 4 describes the prior research that is related to our work. Section 5 discusses the threats to validity. Finally, Section 6 concludes the paper.

2 Methodology

In this section we present our rationale for selecting our studied systems and present our data extraction and analysis approach.

In this paper, we aim to find the unstable logs in a system for easier management of log processing tools. We cloned over 15 projects from git repository locally, and use the ‘grep’ command to recursively search for all log lines within ‘.java’ files in each project. We picked the top four projects with the highest log line count. The four open source projects are: Liferay, Camel, ActiveMQ and CloudStack. All studied systems have extensive system logs and Table 1 presents an overview of the systems.

Table 1: An Overview of all studied Systems

Projects	Liferay	Camel	ActiveMQ	CloudStack
Starting Release	6.1.0-b3	1.6.0	4.1.1	2.1.3
End Release	7.0.0-m3	2.11.3	5.9.0	4.2.0
Total # of Releases	24	43	19	111
Total Added Code	3.9M	505k	261k	1.09M
Total Deleted Code	2.8M	174k	114k	750K
Total Added Log	10.4k	5.1k	4.5k	24k
Total Deleted Log	8.1k	2.4k	2.3k	17k

Liferay¹: Liferay is a free and open source enterprise project written in Java. It provides platform features which are used in development of websites and portals. We select liferay as it is one of the leading platforms for platform development [11] and has growing user base [12]. It also has an extensive issue tracking system in JIRA which helps in categorizing the commits. We study the releases from 6.1.0 to 7.0.0-m3 which cover more than three years of development from 2010 to 2014

¹ <http://www.liferay.com/>




Fig. 1: Flowchart to categorize the different types of log changes that occur

Camel²: Camel is an open source integration platform based on enterprise integration patterns. We select camel as it has an extensive development history and is well logged. We analyze Camel release 1.6 to 2.11.3 which cover more than five years of development from 2009 to 2013.

ActiveMQ³: ActiveMQ is an open source message broker and integration patterns server. We select ActiveMQ as it a messaging oriented middle ware and monitored by wide array of tools [13], making it ideal for our study. We covered ActiveMQ release 4.1.1 to 5.9.0 which cover more than 6 years of development from 2007 to 2013.

CloudStack⁴: Apache CloudStack is open source software designed to deploy and manage large networks of virtual machines, as a highly available, highly scalable Infrastructure as a Service (IaaS) cloud computing platform. We select CloudStack as it as new application and has become one of the four biggest market share currently in cloud domain [14]. We covered CloudStack release 2.1.3 to 4.20 which cover more than 3 year of development from 2010 to 2014.

In the remainder of this section we present our approach for preparing the data to answer our research questions.

2.1 Extracting code evolution

In order to find the stability of logs we have to identify all the ‘Java’ files in our studied systems. To achieve this, we clone the *master* branch of each studied system locally. We use the ‘find’ command to recursively find all the files which end with pattern ‘*.java’. To remove the *Java Test* files, we use ‘grep’ command to filter all files which have ‘test’ or ‘Test’ in their pathname.

After collecting all the Java files from the four studied systems, we use their git repository to obtain all the changes made to the file within the time-frame discussed above. We use the ‘follow’ option to track the file even when they are renamed or relocated within our studied systems. We flatten all the

² <http://camel.apache.org/>

³ <http://activemq.apache.org/>

⁴ <https://cloudstack.apache.org/>

changes made to files to include the changes made in different branches but exclude the final merging commit. Using this approach, we obtain a complete history of each *Java* file.

2.2 Identifying log changes

From the extracted code evolution data for each *Java* file, we identify all the log changes made in the commits. To identify the log statements in the source code, we manually sample some commits from each studied system and identify the logging library used to generate the logs. We find that the studied systems use *Log4j*¹ [19] and *Slf4j*² widely and *logback*³ sparingly. Using this we identify the common method invocations that invoke the logging library. For example in ActiveMQ and Camel a logging library is invoked by method named ‘LOG’ as shown below.

```
LOG.debug("Exception detail", exception);
```

In CloudStack it is usually done through ‘logger’ as follows.

```
_logger.warn("Timed out: " + buildCommandLine(command));
```

As projects can have multiple logging libraries throughout its life-cycle, we use regular expressions to match all the common log invocation patterns (i.e., ‘LOG’, ‘log’, ‘_logger’, ‘LOGGER’, ‘Log’). We count every such invocation of logging library as a log statement.

2.3 Tracking log changes

After identifying all the log changes made to a file across multiple commits, we track each log individually to find out whether it has changed in subsequent revisions. To track the log changes made, we used Levenshtein ratio [15]. We first collect all the logs present in a file at the first commit, which form the initial set of log statements for the file. Every subsequent commit which has changes to a log appears as an added and deleted log statement in *git*.

We consider a pair of added and deleted logs as a modification if they have a Levenshtein ratio of 0.6 or higher. To calculate Levenshtein ratio for a pair of added and deleted logs, we first remove the logging method (e.g., LOG) and compare the remaining text. For example in the logs shown below, we remove the logging method (i.e., LOG) and find that the Levenshtein ratio between the added and deleted pair is 0.86. Hence, this log change is categorized as a log modification.

¹ <http://logging.apache.org/log4j/2.x/>

² <http://www.slf4j.org/>

³ <http://logback.qos.ch/>

```
+ LOG.debug("Call: " +method.getName()+" took "+ callTime +
"ms");
- LOG.debug("Call: " +method.getName()+ " " + callTime);
```

If the added log does not match any log in our initial set, it is considered as a new log addition and is added to the initial set for tracking in future commits. From this we can track how many times a log is changed and how many commits are made between the changes.

2.4 Match JIRA to log changes

Using the commit data, we also track the JIRA issues to extract the developer metrics. To achieve this, we extract the JIRA issue IDs from commit messages and use the JIRA repository to extract the JIRA issue. The JIRA issue contains information about the issue such as, type (i.e., bug, improvement, new-feature, task), priority, resolution time and developer information such as, # of comments and # of developers involved in the JIRA discussion. We use this information along with code and log churn metrics for answering our research questions.

3 Study Results

In this section, we present our study results by answering our research questions. For each question, we discuss the motivation behind it, the approach to answering it and finally the results obtained.

RQ1: How much do logs change over time and why do the changes occur?

Motivation

Research has shown that logs evolve along with the code [16]. When logs are changed, the log processing tools which are dependent on them also have to get updated. This results in costly maintenance effort. To understand the cost, we have to understand how frequent changes to logs are. Hence, we explore the frequency of changes to logs in all our studied systems.

Approach

To find the frequency of log changes, we conduct a quantitative analysis on our studied systems. We use the tracked log data for each studied system as explained in Section 2. From each project, we select a random sample with 95% confidence interval. We follow the same iterative process as in prior research [17] to find how frequently logs change in our studied systems.

When logs are changed, they can be changed in several ways. To understand the different types of log changes we perform a manual analysis on the

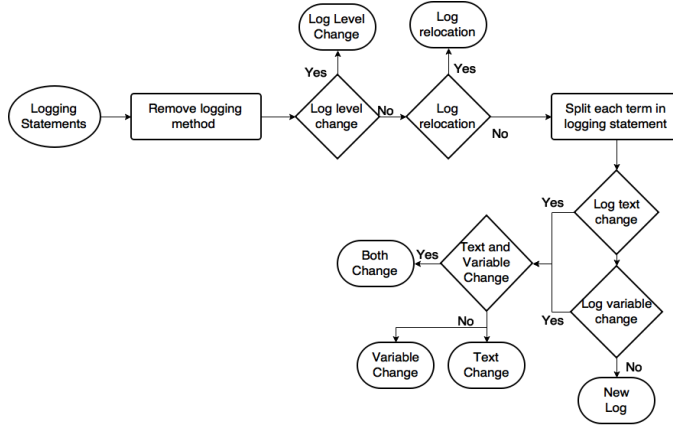


Fig. 2: Flowchart to categorize the different types of log changes that occur

changed logs. We select a random sample from each project such that the sample achieves 95% interval. After identifying all the types of log changes we automate the process using *Levenshtein ratio* [15]. Figure 2 highlights the process of categorizing the log changes. For example, consider the logs shown below.

```
+ LOG.info("starting HBase HsHA Thrift server on " + Integer.toString(listenPort));
```

```
- LOG.info("starting HBase " + implType.simpleClassName() + "server on " + Integer.toString(listenPort));
```

To identify the type of log change we first remove the logging method (i.e, LOG) and the log level (i.e, info) from the logs. We then compute the *Levenshtein ratio* between each term within the parentheses. In the above case we find that '+ Integer.toString(listenPort)' has *Levenshtein ratio* of 1, implying they are identical and the *Levenshtein ratio* between 'starting HBase HsHA Thrift server on' and 'starting HBase' is 0.56. This suggests there is some similarity between the two strings. Using this information we categorize logs into different types.

Results

We find that developers change 20-80% of the logs across our studied systems as shown in Table 2. Based on frequency of changes, we categorize logs into 3 categories namely: a) Frequently Changed, b) Changed and c) Never Changed. If a log is changed more than three times it is categorized as 'Frequently Changed'. If it is changed fewer than three times it is categorized under 'changed' and if there is no changes made it is categorized under 'Never Changed'. We see that 'Never Changed' is majority only in Liferay and Camel but is lesser than 'Changed' in ActiveMQ and CloudStack, as

Table 2: Distribution of log changes in different projects

Projects	Never Changed (%)	Changed (%)	Frequently Changed (%)
Life Ray	78.67	19.66	1.66
Camel	55.43	37.32	7.25
ActiveMq	34.78	62.02	3.20
CloudStack	19.68	68.61	11.71

Table 3: Distribution of log changes in different projects

Projects	Life Ray	Cloud Stack	Active-MQ	Camel
Log Relocation (%)	73.5	78.85	63.27	49.72
Log Text Change (%)	20.11	6.37	5.97	7.59
Log Variable Change (%)	3.10	7.21	6.91	8.29
Change of Log Level (%)	0.87	1.15	1.76	3.79
Text and Variable Change (%)	2.33	6.39	22.0	30.59

seen in Table 2. This may be because Camel and Liferay rely less on logs as their middle-ware/application software, whereas ActiveMQ and CloudStack are service software.

When developers change logs there can be many types of changes, i.e., changes to log context, verbosity levels or logged variables. From our manual analysis, we identify five types of log changes. Table 3 shows their distributions. When there is overlapping of the different types of log changes, we categorize them as newly added log and track changes made to it.

1. **Log relocation:** The log is kept intact but moved to different location in the file because of context changes(code around the log is changed).
2. **Text change:** The text (i.e., static content) of log is changed.
3. **Variable change:** One or more variables in the log are changed (added, deleted or modified).
4. **Change of log level:** The verbosity level of a log is changed.
5. **Text and variable change:** Both text and variables in the logs are changed. This is generally done when developers provide more context information, i.e, text and add/modify the relevant variables in a log.

Log relocation occurs more than other types of log changes. We find that in =all the studied systems log relocation occurs more than 40%. As log relocations have no changes to the text or variables in logs, their impact on log processing tools is limited. Hence we exclude log relocation changes from our datasets.

45% of the logs are changed atleast once in three of our studied systems. We find that over 50% of these log changes are due to log relocation. The remaining changes are to the text (i.e., static content), variable and log level. This suggests that developers change logs extensively throughout the lifecycle of a software and developers might dedicate significant amount of time to update and maintain logs


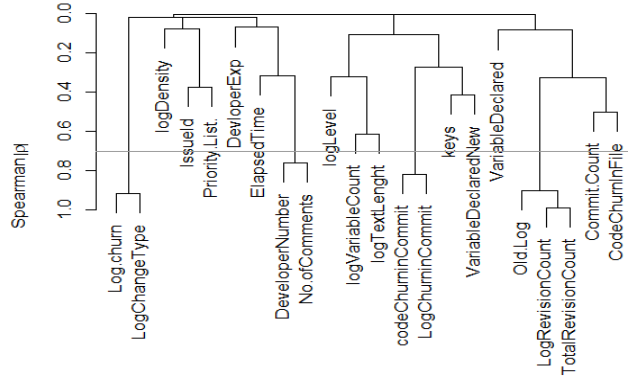


Fig. 3: Overview of Model construction(C) and Evaluation

Fig. 4: Hierarchical clustering of variables according to Spearman's ρ in ActiveMQ

RQ2: Can code and developer related metrics help in explaining the stability of logs?

In RQ1, we find that 20-80% of logs are changed in our studied systems. This affects the log processing tools which run on these studied systems, making developers spend more time on maintenance of those tools. Hence, there is a need to identify these non-stable logs to simplify the job of developers. It can also benefit log processing tool developers to develop more robust applications making them more stable.

Approach

To find the stability of logs in our studied systems, we extract code and log churn metrics from the Git repository and developer metrics from JIRA. We use code churn metrics because prior work has linked logs to development knowledge and issue reports [17].

Table 4 lists all the metrics we collect for each dimension. We define each metric and the rationale behind the choice of each metric. We collect the code and log churn metrics when a log is added into the studied system and track the log statement. If it is changed in subsequent revisions we flag the log as changed and record it.

Table 4: Taxonomy of considered metrics for model construction

Dimension	Metrics	Values	Definition (d) – Rationale (r)
Code Metrics	Log revision count	Numeric	d: The number of commits prior which had log statement changes.
			r: This helps to identify if the file is prone to log statement changes.
	New File	Boolean (0 -1)	d: Check if the log is added in a new file (i.e., newly committed)
			r: This helps to identify which logs were added later in subsequent commits.
	Total Revision Count	Numeric	d: Total number of commits made to the file before the log is added. This value is 0 for logs added in the initial commit but not for logs added overtime.
			r: This helps to find out if the file is changed heavily which can result in log changes [18]
	Code churn in commit	Numeric	d: The code churn of the commit in which a log is added.
			r: Log changes are correlated to code churn in files [19]
	Variables declared	Numeric	d: The number of variables which are declared before the log statement. (we limit to 20 lines before log statement)
			r: When new variables are declared, developers may log the new variables to obtain more information [18]
Log Metrics	Log Context	Categorical	d: Identify the block in which a log is added. (i.e., 'if', 'if=else', 'try-catch', 'exception', 'throw', 'new function')
			r: Prior research find that logs are mostly used in assertion checks, logical branching, return value checking, assertion checking [21]
	Log Level	Categorical	d: Identify the log level (verbosity) of the added log (i.e., 'info', 'error', 'warn', 'debug', 'trace' and 'trace')
			r: Developers spend significant amount of time in adjusting the verbosity of logs [18]
	Log variable count	Numerical	d: Number of variables logged.
			r: Over 62% of logs add new variables [18]. Hence fewer variables in the initial log statement might result in addition later.
	Log text length	Numerical	d: Number of text phrases logged (i.e., we count all text present between a pair of quotes as one phrase)
			r: Over 45% of logs have modifications to static context [18]. Logs with fewer phrases might be subject to changes later to provide better explanation
	Log density	Numerical	d: Ratio of number of log lines to the source code lines in the file.
			r: Research has found that there is on average one log line per 30 lines of code [18]. If it is less it suggests there may be additions in later commits.

Dimension	Metrics	Values	Definition (d) – Rationale (r)
Developer Metrics	Resolution time	Numerical	d: The time it takes for the issue to get fixed. It is defined as the time it takes since an issue is opened until closed.
			r: More resolution time might suggest a more complex fix with more code churn resulting in more log churn.
	Number of developers involved	Numeric	d: Total number of unique developers who comment on the issue report on JIRA
			r: Components with many unique authors likely lack strong ownership, which in turn may lead to more defects [22] and change logs [19].
	Number of comments	Numeric	d: Total number of discussion posts on the issue.
			r: Number of comments is correlated to the resolution time of issue reports [23]. More comments may also indicate the issue is more complex resulting in more code churn and log changes.
	Developer experience	Numeric	d: The number of commits the developer has made prior to this commit.
			r: Research has shown that experienced developers might take up more complex issues [24] and therefore may leverage logs more [19].
	Issue type	Categorical	d: Identify the type of issue, i.e., ‘Bug’, ‘Improvement’, ‘Task’, ‘New Feature’, ‘Sub-Task’, ‘Test’
			r: Some issue types might have higher code churn than others (example: Bug and New features might have more code churn when compared to Sub-Tasks) and are committed faster.
	Priority type	Categorical	d: Identify the priority of the issue i.e., ‘Critical’, ‘Blocker’, ‘Major’, ‘Minor’ and ‘Trivial’
			r: Research has shown that priority of issue affects resolution time of bug fixes [25]. A Higher priority indicates the issue will be fixed faster with log changes.

We adopt the statistical tool R to model our data and use the ‘Random-Forest’ package to generate the random forests. The overview of the process is shown in Figure 3 and explained below.

(C1 - Correlation analysis)

Correlation analysis is necessary to remove the highly correlated metrics from our dataset. Collinearity between metrics can effect the performance of a model because, small changes in one metric can affect the values of other metrics, causing wide swings in predicted parameter.

We use Spearman rank [26] correlation instead of Pearson [27] because Spearman is resilient to data that is not normally distributed. We use the function ‘varclus’ in R to perform the correlation analysis.

Figure 4 shows the hierarchically clustered Spearman ρ values in ActiveMQ project. The solid horizontal lines indicate the correlation value of the two metrics that are connected by the vertical branches that descend from it. We exclude one metric from the sub-hierarchies which have correlation $|\rho| > 0.7$. The gray line indicates our cutoff value ($|\rho| = 0.7$). We use cutoff value of ($|\rho| = 0.7$) as used by prior research [28] to remove the correlated metrics before

	Predicted log is changed	Predicted log is not-changed
Actual log is changed	TP (True Positive)	FN (False Negative)
Actual log is not-Changed	FP (False Positive)	TN (True Negative)

Table 5: Confusion Matrix

building our model.

(C2 - Random forest generation)

After we eliminate the correlated metrics from our datasets, we construct the random forest model. A random forest is a collection of largely uncorrelated decision trees in which the results of all trees are combined to form a generalized predictor [29].

Given a dataset of n logs for training, $D = \{(X_1Y_1, \dots, (X_nY_n))\}$ where $X_i, i = 1 \dots n$, is a vector of descriptors (i.e, explanatory metrics) and Y_i is the flag which indicates whether a log is changed or not, the algorithm follows these steps.

- From the training set of n logs, a random sample of n components is selected with replacement (i.e., bootstrap sample) [28].
We use the ‘boot’ function from the ‘boot’ package in R to generate bootstrap samples. The boot function works by generating a set of random indices, with replacement from the integers $1:n(\text{data})$.
- From the sampled selection, a tree is grown with one difference compared to normal decision trees. The random forest is split at each node using the best among subset of predictors, rather than all. This counterintuitive strategy helps to make the random forests more robust against over-fitting [30].
We use the ‘randomForest’ function from the ‘randomForest’ package in R to generate the random forest model.
- The above steps are repeated until M such models are grown.
- Predict new data by aggregating the prediction of the M models generated.

(C3 - Model evaluation)

After we build the random forest model, we evaluate the performance of our model using precision, recall, F-measure and Brier Score. All these measures are functions of the confusion matrix as shown in Table 5 and are explained below.

Precision (P) measure the correctness of our model in predicting which log will undergo a change in the future. It is defined as the number of logs which were accurately predicted as changed over all logs predicted to have changed as explained in Equation 1.

$$P = \frac{TP}{TP + FP} \quad (1)$$

Recall (R) measure the completeness of our model. A model is said to be complete if the model can predict all the logs which will get changed in our

dataset. It is defined as the number of logs which were accurately predicted as changed over all logs which actually change as explained in Equation 2.

$$R = \frac{TP}{TP + FN} \quad (2)$$

F-Measure is the harmonic mean of precision and recall, combining the inversely related measure into a single descriptive statistic as shown in Equation 3 [31].

$$F = \frac{2 \times P \times R}{P + R} \quad (3)$$

Brier Score (BS) is a measure of the accuracy of the predictions in our model. It is the mean squared error of the probability forecasts [32]. The most common formulation of Brier Score is shown in Equation 4, where f_t is the probability that was predicted, o_t is the actual outcome of the event at the instance t (0 if log is not changed and 1 if it is changed), and N is the number of forecasting instances.

If the predicted value is 70% and the log is changed, the Brier Score is 0.09 (lower the Brier score, the more likely the event will occur). It reaches 0.25 for random guessing (i.e., predicted value is 50%) and it is 0 when predicted value is 100% and the log is not changed.

$$BS = \frac{1}{N} \sum_{t=1}^N (f_t - o_t)^2 \quad (4)$$

The model performance measure, described previously only provide insight into how the random forest models fit the observed dataset, but it may overestimate the performance of the model if it is over-fit. We take performance overestimation (i.e., *optimism*) into account by subtracting the bootstrap averaged optimism [28] from the initial performance measure. The *optimism* of the performance measures are calculated as follows:

1. From the original dataset with n records, select a bootstrap sample with n components with replacement.
2. Build random forest using the bootstrap sample.
3. Apply the model built from bootstrap data on both the bootstrap and original data sample, calculating precision, recall, F-measure and Brier score for both data samples.
4. Calculate *Optimism* by subtracting the performance measures of the bootstrap sample against the original sample.

The above process is repeated 1,000 times and the average (mean) *optimism* is calculated. Finally, we calculate *optimism-reduced* performance measures for precision, recall, F-measure and Brier score by subtracting the averaged optimism of each measure, from their corresponding original measure. The smaller the optimism values, the more stable that the original model fit is.

(C4 - Importance of each metric in relation to log stability) To find the importance of each metric in a random forest model, we use a permutation

Table 6: The optimism reduced performance measures of the four projects

Projects	Life Ray	ActiveMQ	Camel	CloudStack
Precision	0.91	0.90	0.93	0.89
Recall	0.76	0.67	0.89	0.92
F-Measure	0.83	0.77	0.91	0.90
Brier Score	0.06	0.04	0.06	0.07

test. In this, for every tree grown in the forest, we use the test dataset (out-of-bag) to predict the accuracy of the model [33]. Then, the values of the X_i th metric of which we want to find importance for, are randomly permuted in the test dataset and the accuracy of the model is recomputed. The decrease in accuracy as a result of this permutation is averaged over all trees, and is used as a measure of the importance of metric X_i th in the random forest.

We use the ‘importance’ function defined in ‘RandomForest’ package of R, to calculate the importance of each metric. We call the ‘importance’ function each time during the bootstrapping process to obtain 1000 importance scores for each metric in our dataset.

We use the **Scott-Knott** clustering to group the metric based on their means [34,35]. This is done to group metrics which are strong predictors of likelihood of log change. The SK algorithm uses the hierarchical clustering approach to divide the metrics and uses the likelihood ratio test to judge the difference between the groups. This assures the means of metrics within a group not to be statistically significantly different. We use the ‘SK’ function in the ‘ScottKnott’ package of R and set significance threshold of 0.05 to cluster the metrics into different groups.

Results

The random forest classifier achieves a precision of 0.89-0.93 and recall of 0.76-0.92 for our studied systems. Table 6 shows the optimism-reduced values of *precision*, *recall*, *F-measure* and *Brier score* for each project. We find the recall for random forest classifier in Liferay is not as good as the other projects. This may be because Liferay has the lowest percentage of logs which are changed at 20%, compared to 45-70% in the other projects. Because, of the lower percentage of logs which are changed, the random forest classifier has fewer nodes (trees) and likelihood of false negatives is higher.

The random forest classifier outperforms random guessing. The classifier achieves Brier scores between 0.04 and 0.07 across all projects. As the Brier score measure the total difference between the event and the forecast probability of the event, a perfect classifier will have Brier score of 0 and prefect misfit classifier will have Brier score of 1(predicts probability of log change when log not changed). This means lower the Brier score value, better fit our random forest classifier.

We find that ‘Developer Experience’ is the most important factor for developer metrics. From Code Metrics we find that ‘SLOC’, ‘Code Churn in Commit’ and ‘# of comments’ are the important factors and from log metrics

Table 7: The important values of the metrics (top 10), divided into homogeneous groups by Scott-Knott clustering. The ‘+’ and ‘-’ signs signify positive and negative correlation of the metric on log stability.

ActiveMQ			CloudStack		
Groups	Factors	Importance	Groups	Factors	Importance
1	Developer experience	0.144 -	1	Log revision count	0.116 +
2	Log text length	0.103 +	2	SLOC	0.078 +
3	Developer number	0.098 +	3	Developer Experience	0.073 -
4	SLOC	0.097 +	4	Priority list	0.067 +
5	Log density	0.082 +	5	# of comments	0.064 -
6	Priority list	0.076 -	6	Log density	0.063 -
7	Resolution time	0.070 +	7	Code churn in commit	0.056 +
8	# of comments	0.070 +	8	Variable declared	0.042 +
	Code churn in commit	0.067 +	9	Log text length	0.039 +
	Log revision count	0.067 +	10	Log context	0.037 -
Life Ray			Camel		
Groups	Factors	Importance	Groups	Factors	Importance
1	Log context	0.276 -	1	Log text length	0.157 +
2	Log density	0.270 +	2	Developer experience	0.147 -
3	Resolution time	0.162 -	3	Log variable count	0.086 +
4	Developer experience	0.156 -	4	Log density	0.082 +
5	SLOC	0.143 +	5	Log revision count	0.081 +
6	Log level	0.097 -		# of comments	0.081 -
7	Log variable count	0.090 +	6	Code churn in commit	0.079 -
	# of comments	0.089 -	7	Resolution time	0.078 -
8	Priority list	0.079 +	8	SLOC	0.077 +
9	Log text length	0.0 +	9	Log level	0.065 -

we find that ‘log density’ is the important factor. Table 7 shows the important factors of the predictor metrics, divided into distinct groups. If two or more metrics have similar importance values they are grouped together. The ‘+’ or ‘-’ sign following the predictor metric indicates if the metric has positive or negative affect on the stability of logs.

We find that *log density*, *log text length*, *log revision count* are important metrics in our studied systems as shown in Table 7. We find *log text length* has a positive correlation, implying more text in a log, higher the likelihood that log will be changed. This may be because there are inconsistencies between logs and the actual needed information intended as shown by prior research [18] and developers have to update logs to resolve the inconsistencies.

Log revision count has positive correlation on log change, across all studied systems. This is because, if a log in a file is changed often, there is higher chance of it getting changed again, than logs in file with no prior changes [36].

We find that *SLOC* and *# of comments* are strong predictors of log change, from *Code Churn* dimension across all projects. *SLOC* (source lines of code) has positive correlation suggesting that logs in larger files have higher likelihood of getting changed. We find that *# of comments* has negative correlation which suggests that when developers provide information through comments, logs around those comments are less likely to get changed.

We find that *developer experience* is a strong predictor of log change in all our models and it has negative correlation across all projects. This implies that logs are more likely to be changed by less experienced developers. We find that

other metrics from Developer dimension are not consistent among the studied systems. This may be because each project might have different philosophy of development, for example we find that *resolution time* has negative affect in Liferay and Camel but has positive affect in ActiveMQ. This suggests that logs are more likely to change in ActiveMQ when the resolution time of issue increases, but less likely in Liferay and Camel.

Our Random Forest classifier achieves a precision of 0.89-0.93 and a recall of 0.76-0.92 across all studied systems. We find developer experience, SLOC, # of comments and log density to be strong predictors of log change in our studied systems

4 Related Work

In this section, we present prior research that performs log analysis on large software systems and tools developed to assist in logging.

4.1 Log Analysis

Prior work leverages logs for testing and detecting anomalies in large scale systems. Shang *et al.* [16] propose an approach to leverage logs in verifying the deployment of Big Data Analytics applications. Their approach analyzes logs in order to find differences between running in a small testing environment and a large field environment. Lou *et al.* [2] propose an approach to use the variable values printed in logs to detect anomalies in large systems. Based on the variable values in logs, their approach creates invariants (e.g., equations). Any new logs that violates the invariant are considered to be a sign of anomalies. Fu *et al.* [3] built a Finite State Automaton (FSA) using unstructured logs to detect performance bugs in distributed systems.

Xu *et al.* [1] link output logs to logs in source code to recover the text and the variable parts of logs in source code. They applied Principal Component Analysis (PCA) to detect system anomalies. Tan *et al.* [37] propose a tool named SALSA, which constructs state-machines from logs. The state-machines are further used to detect anomalies in distributed computing platforms. Jiang *et al.* [38] study the leverage of logs in troubleshooting issues from storage systems. They find that logs assist in a faster resolution of issues in storage systems. Beschastnikh *et al.* [39] designed automated tools that infers execution models from logs. These models can be used by developers to understand the behaviours of concurrent systems. Moreover, the models also assist in verifying the correctness of the system and fixing bugs.

To assist in fixing bugs using logs, Yuan *et al.* [40] propose an approach to automatically infer the failure scenarios when a log is printed during a failed run of a system.

Jiang *et al.* [41, 10, 42, 43] proposed log analysis approaches to assist in automatically verifying results from load tests. Their log analysis approaches

first automatically abstract logs into system events [41]. Based on the such events, they identified both functional anomalies [10] and performance degradations [42] in load test results. In addition, they proposed an approach that leverage logs to reduce the load test that are performed in user environment [43].

The extensive prior research of log analysis shows that logs are leveraged for different purposes and changes to logs can effect performance of log analysis tools. This motivates our paper to study how much logs are changed by developers. As a first step, we study how many logs are stable and how many undergo changes across commits. Our findings show that more 20-80% of the logs are changed at-least once and 3-11% of the logs are changed frequently.

4.2 Empirical studies and tools developed on logs

Prior research performs an empirical study on the characteristics of logs. Yuan *et al.* [18] studies the logging characteristics in four open source systems. They find that over 33% of all log changes are after thoughts and logs are changed 1.8 times more regular entire code. Fu *et al.* [21] performed an empirical study on where developers put logs. They find that logs are used for assertion checks, return value checks, exceptions, logic-branching and observing key points. The results of the analysis were evaluated by professionals from the industry and a F-score of over 95% was achieved.

Shang *et al.* [44] signify the fact that there is a gap between operators and developers of software systems, especially in the leverage of logs. They performed an empirical study on the evolution of both static logs and logs outputted during run time [19,45]. They find that logs are co-evolving with software systems. However, logs are often modified by developers without considering the needs of operators. Furthermore, Shang *et al.* [17] find that understanding logs is challenging. They examine user mailing lists from three large open-source projects and find that users of these systems have various issues in understanding logs outputted by the system. Shang *et al.* propose to leverage different types of development knowledge, such as issue reports, to assist in understanding logs.

The most related research by Yuan *et al.* [5] shows that logs need to be improved by providing additional information. Their tool named *Log Enhancer* can automatically provide additional control and data flow parameters into the logs thereby improving the logs. Our paper focuses more on informing developers which logs are more likely to get changed and identifying which factors explain the change of logs.

5 Threats to Validity

In this section, we present the threats to the validity to our findings.

External Validity

Our case study is performed on Liferay, ActiveMQ, Camel and CloudStack. Though these studied systems have years of history and large user bases, these systems are all Java based platform softwares. Software in other domains may not rely on logs and may not use log processing tools. Our projects are all open source and we do not verify the results on any commercial platform systems. More case studies on other domains and commercial platforms, with other programming languages are needed to see whether our findings can be generalized.

Construct Validity

Our heuristics to extract logging source code may not be able to extract every log in the source code. Even though the studied systems-leverage logging libraries to generate logs at runtime, there still exist user-defined logs. By manually examining the source code, we believe that we extract most of the logs. Evaluation on the coverage of our extracted logs can address this threat.

We use keywords to identify bug fixing commits when the JIRA issue ID is not included in the commit messages. Although such keywords are used extensively in prior research [19], we may still fail to identify bug fixing commits or branching and merging commits.

We use Levenshtein ratio and choose a threshold to identify modifications to logs. However, such threshold may not accurately identify modifications to logs. Further sensitivity analysis on such threshold is needed to better understand the impact of the threshold to our findings.

Our models are incomplete, i.e., we have not measured all potential dimensions that impact logging stability. However, we feel that our metrics are good and additional new dimensions can compliment our dimension to improve the explanation.

Internal Validity

Our study is based on the data obtained from Git and JIRA for all the studied systems. The quality of the data contained in the repositories can impact the internal validity of our study.

Our analysis of the relationship between metrics that are important factors in predicting the stability of logs cannot claim causal effects, as we are investigating correlation and not causation. The important factors from our random forest models only indicate that there exists a relationship which should be studied in depth in future studies.

6 Conclusion

Logs are snippets of code, introduced by experts to record valuable information. This information is used by a plethora of log processing tools to assist in software testing, monitoring performance and system state comprehension. These log processing tools are completely dependent on the logs and hence are impacted when logs are changed.

In this paper we study the stability of logs, which affects the performance of the processing tools which utilize them. We build a Random Forest classifier to predict which logs are more likely to change in the future using data from three dimensions namely: *code*, *log* and *developers*. We collect this data when logs are added in the source code and we predict which log will change in the future using this data. The highlights of our study are:

- We find that 20-80 % of logs are changed at-least once.
- Our Random forest classifier achieves a precision of 89-93% and recall of 76-92%.
- We find a negative correlation between developer experience and log changes from the developer dimension. We find that the number of comments in source code, also has negative correlation towards log changes.

Our findings highlight that we can predict which logs have a higher likelihood of getting changed when they are introduced. This information can be used by system administrators and practitioners, to identify logs which have a higher chance of affecting the log processing tools and prevent failure of these tools.

References

1. W. Xu, L. Huang, A. Fox, D. Patterson, and M. I. Jordan, “Detecting large-scale system problems by mining console logs,” in *ACM SIGOPS ’09: Proceedings of the 22nd Symposium on Operating Systems Principles*, New York, NY, USA, pp. 117–132.
2. J.-G. Lou, Q. Fu, S. Yang, Y. Xu, and J. Li, “Mining invariants from console logs for system problem detection,” in *USENIX Annual Technical Conference*, 2010.
3. Q. F. J. L. Y. Wang and J. Li., “Execution anomaly detection in distributed systems through unstructured log analysis,” in *ICDM 09: In Proceedings of 9th IEEE International Conference on Data Mining*.
4. H. Malik, H. Hemmati, and A. Hassan, “Automatic detection of performance deviations in the load testing of large scale systems,” in *Software Engineering (ICSE), 2013 35th International Conference on*, May 2013, pp. 1012–1021.
5. D. Yuan, J. Zheng, S. Park, Y. Zhou, and S. Savage, “Improving software diagnosability via log enhancement,” *ACM Transactions on Computer Systems*, vol. 30, pp. 4:1–4:28, Feb. 2012.
6. W. Shang, Z. M. Jiang, B. Adams, A. E. Hassan, M. W. Godfrey, M. Nasser, and P. Flora, “An exploratory study of the evolution of communicated information about the execution of large software systems,” *Journal of Software: Evolution and Process*, vol. 26, no. 1, pp. 3–26, 2014.
7. A. E. Hassan, D. J. Martin, P. Flora, P. Mansfield, and D. Dietz, “An industrial case study of customizing operational profiles using log compression,” in *Software Engineering, 2008. ICSE’08. ACM/IEEE 30th International Conference on*. IEEE, 2008, pp. 713–723.
8. M. Nagappan, K. Wu, M. Vouk *et al.*, “Efficiently extracting operational profiles from execution logs using suffix arrays,” in *Software Reliability Engineering, 2009. ISSRE’09. 20th International Symposium on*. IEEE, 2009, pp. 41–50.
9. L. Bitincka, A. Ganapathi, S. Sorkin, S. Zhang *et al.*, “Optimizing data analysis with a semi-structured time series database,” in *Proceedings of the 2010 workshop on Managing systems via log analysis and machine learning techniques*. USENIX Association, 2010, pp. 7–7.

10. Z. M. Jiang, A. Hassan, G. Hamann, and P. Flora, "Automatic identification of load testing problems," in *ICSM '08: Proceedings of the IEEE International Conference on Software Maintenance*, Sept 2008, pp. 307–316.
11. "Liferay gartner report - <https://www.liferay.com/web/bryan.cheung/blog/-/blogs/liferay-s-thoughts-on-the-2014-gartner-magic-quadrant-for-horizontal-portals>."
12. "Liferay user report - <http://trends.builtwith.com/cms/liferay>."
13. "Activemq monitors - <http://activemq.apache.org/how-can-i-monitor-activemq.html>."
14. "Cloudstack - <http://www.datacenterknowledge.com/archives/2015/02/23/openstack-vs-cloudstack-the-platforms-and-the-cloud-apis/>."
15. M. Mednis and M. K. Aurich, "Application of string similarity ratio and edit distance in automatic metabolite reconciliation comparing reconstructions and models," *Biosystems and Information technology*, vol. 1, no. 1, pp. 14–18, 2012.
16. W. Shang, Z. M. Jiang, H. Hemmati, B. Adams, A. E. Hassan, and P. Martin, "Assisting developers of big data analytics applications when deploying on hadoop clouds," in *ICSE'13: Proceedings of the 2013 International Conference on Software Engineering*, ser. ICSE '13. Piscataway, NJ, USA: IEEE Press, 2013, pp. 402–411.
17. W. Shang, M. Nagappan, A. E. Hassan, and Z. M. Jiang, "Understanding log lines using development knowledge," in *ICSME '14: Proceedings of the International Conference on Software Maintenance and Evolution*,. IEEE, 2014, pp. 21–30.
18. D. Yuan, S. Park, and Y. Zhou, "Characterizing logging practices in open-source software," in *ICSE '12: Proceedings of the 34th International Conference on Software Engineering*. IEEE Press, 2012, pp. 102–112.
19. W. Shang, M. Nagappan, and A. E. Hassan, "Studying the relationship between logging characteristics and the code quality of platform software," *Empirical Software Engineering*, vol. 20, no. 1, pp. 1–27, 2015.
20. D. Zhang, K. El Emam, H. Liu *et al.*, "An investigation into the functional form of the size-defect relationship for software modules," *Software Engineering, IEEE Transactions on*, vol. 35, no. 2, pp. 293–304, 2009.
21. Q. Fu, J. Zhu, W. Hu, J.-G. L. R. Ding, Q. Lin, D. Zhang, and T. Xie, "Where do developers log? an empirical study on logging practices in industry," in *ICSE Companion '14: Proceedings of the 36th International Conference on Software Engineering*,. pp. Pages 24–33.
22. S. McIntosh, Y. Kamei, B. Adams, and A. E. Hassan, "The impact of code review coverage and code review participation on software quality: A case study of the qt, vtk, and itk projects," in *Proceedings of the 11th Working Conference on Mining Software Repositories*. ACM, 2014, pp. 192–201.
23. E. Giger, M. Pinzger, and H. Gall, "Predicting the fix time of bugs," in *Proceedings of the 2nd International Workshop on Recommendation Systems for Software Engineering*. ACM, 2010, pp. 52–56.
24. F. Rahman and P. Devanbu, "Ownership, experience and defects: a fine-grained study of authorship," in *Proceedings of the 33rd International Conference on Software Engineering*. ACM, 2011, pp. 491–500.
25. L. Marks, Y. Zou, and A. E. Hassan, "Studying the fix-time for bugs in large open source projects," in *Proceedings of the 7th International Conference on Predictive Models in Software Engineering*. ACM, 2011, p. 11.
26. J. H. Zar, *Spearman Rank Correlation*. John Wiley & Sons, Ltd, 2005. [Online]. Available: <http://dx.doi.org/10.1002/0470011815.b2a15150>
27. R. J. Serfling, *Approximation theorems of mathematical statistics*. John Wiley & Sons, 2009, vol. 162.
28. S. McIntosh, Y. Kamei, B. Adams, and A. E. Hassan, "An empirical study of the impact of modern code review practices on software quality," *Empirical Software Engineering*, p. To appear, 2015.
29. J. Albert, E. Aliu, H. Anderhub, P. Antoranz, A. Armada, M. Asensio, C. Baixeras, J. Barrio, H. Bartko, D. Bastieri, J. Becker, W. Bednarek, K. Berger, C. Bigongiari, A. Biland, R. Bock, P. Bordas, V. Bosch-Ramon, T. Bretz, I. Britvitch, M. Camara, E. Carmona, A. Chilingarian, S. Ciprini, J. Coarasa, S. Commichau, J. Contreras, J. Cortina, M. Costado, V. Curtef, V. Danielyan, F. Dazzi, A. D. Angelis, C. Delgado, R. de los Reyes, B. D. Lotto, E. Domingo-Santamara, D. Dorner, M. Doro, M. Errando,

- M. Fagiolini, D. Ferenc, E. Fernandez, R. Firpo, J. Flix, M. Fonseca, L. Font, M. Fuchs, N. Galante, R. Garca-Lpez, M. Garczarczyk, M. Gaug, M. Giller, F. Goebel, D. Hakobyan, M. Hayashida, T. Hengstebeck, A. Herrero, D. Hhne, J. Hose, S. Huber, C. Hsu, P. Jacon, T. Jogler, R. Kosyra, D. Kranich, R. Kritzer, A. Laille, E. Lindfors, S. Lombardi, F. Longo, J. Lpez, M. Lpez, E. Lorenz, P. Majumdar, G. Maneva, K. Mannheim, M. Mariotti, M. Martnez, D. Mazin, C. Merck, M. Meucci, M. Meyer, J. Miranda, R. Mirzoyan, S. Mizobuchi, A. Moralejo, D. Nieto, K. Nilsson, J. Ninkovic, E. Oa-Wilhelmi, N. Otte, I. Oya, M. Panniello, R. Paoletti, J. Paredes, M. Pasanen, D. Pascoli, F. Pauss, R. Pegna, M. Persic, L. Peruzzo, A. Piccioli, N. Puchades, E. Prandini, A. Raymers, W. Rhode, M. Rib, J. Rico, M. Rissi, A. Robert, S. Rgamer, A. Saggion, T. Saito, A. Snchez, P. Sartori, V. Scalzotto, V. Scapin, R. Schmitt, T. Schweizer, M. Shayduk, K. Shinozaki, S. Shore, N. Sidro, A. Sillanp, D. Sobczynska, F. Spanier, A. Stamerra, L. Stark, L. Takalo, P. Temnikov, D. Tescaro, M. Teshima, D. Torres, N. Turini, H. Vankov, A. Venturini, V. Vitale, R. Wagner, T. Wibig, W. Wittek, F. Zandanel, R. Zanin, and J. Zapatero, "Implementation of the random forest method for the imaging atmospheric cherenkov telescope {MAGIC}," *Nuclear Instruments and Methods in Physics Research Section A: Accelerators, Spectrometers, Detectors and Associated Equipment*, vol. 588, no. 3, pp. 424 – 432, 2008. [Online]. Available: <http://www.sciencedirect.com/science/article/pii/S0168900207024059>
30. L. Breiman, "Bagging predictors," *Machine learning*, vol. 24, no. 2, pp. 123–140, 1996.
 31. G. Hripcsak and A. S. Rothschild, "Agreement, the f-measure, and reliability in information retrieval," *Journal of the American Medical Informatics Association*, vol. 12, no. 3, pp. 296–298, 2005.
 32. D. S. Wilks, *Statistical methods in the atmospheric sciences*. Academic press, 2011, vol. 100.
 33. V. Svetnik, A. Liaw, C. Tong, J. C. Culberson, R. P. Sheridan, and B. P. Feuston, "Random forest: a classification and regression tool for compound classification and qsar modeling," *Journal of chemical information and computer sciences*, vol. 43, no. 6, pp. 1947–1958, 2003.
 34. A. Scott and M. Knott, "A cluster analysis method for grouping means in the analysis of variance," *Biometrics*, pp. 507–512, 1974.
 35. E. G. Jelihovschi, J. C. Faria, and I. B. Allaman, "Scottknott: a package for performing the scott-knott clustering algorithm in r," *TEMA (São Carlos)*, vol. 15, no. 1, pp. 3–17, 2014.
 36. A. E. Hassan, "The road ahead for mining software repositories," in *Frontiers of Software Maintenance, 2008. FoSM 2008*. IEEE, 2008, pp. 48–57.
 37. J. Tan, X. Pan, S. Kavulya, R. Gandhi, and P. Narasimhan, "Salsa: Analyzing logs as state machines," in *WASL'08: Proceedings of the 1st USENIX Conference on Analysis of System Logs*. USENIX Association, 2008, pp. 6–6.
 38. W. Jiang, C. Hu, S. Pasupathy, A. Kanevsky, Z. Li, and Y. Zhou, "Understanding customer problem troubleshooting from storage system logs," in *FAST '09: Proceedings of the 7th Conference on File and Storage Technologies*. Berkeley, CA, USA: USENIX Association, 2009, pp. 43–56.
 39. I. Beschastnikh, Y. Brun, S. Schneider, M. Sloan, and M. D. Ernst, "Leveraging existing instrumentation to automatically infer invariant-constrained models," in *Proceedings of the 19th ACM SIGSOFT Symposium and the 13th European Conference on Foundations of Software Engineering*, ser. ESEC/FSE '11. New York, NY, USA: ACM, 2011, pp. 267–277.
 40. D. Yuan, H. Mai, W. Xiong, L. Tan, Y. Zhou, and S. Pasupathy, "Sherlog: Error diagnosis by connecting clues from run-time logs," in *ASPLOS '10: Proceedings of the 15th Edition of ASPLOS on Architectural Support for Programming Languages and Operating Systems*. New York, NY, USA: ACM, 2010, pp. 143–154.
 41. Z. M. Jiang, A. E. Hassan, G. Hamann, and P. Flora, "An automated approach for abstracting execution logs to execution events," *Journal of Software Maintenance Evolution*, vol. 20, no. 4, pp. 249–267, Jul. 2008.
 42. —, "Automated performance analysis of load tests," in *ICSM '09: Proceedings of the IEEE International Conference on Software Maintenance*. IEEE, 2009, pp. 125–134.

43. Z. M. Jiang, A. Avritzer, E. Shihab, A. E. Hassan, and P. Flora, "An industrial case study on speeding up user acceptance testing by mining execution logs," in *SSIRI '10: Proceedings of the 4th International Conference on Secure Software Integration and Reliability Improvement*. Washington, DC, USA: IEEE Computer Society, pp. 131–140, 2010.
44. W. Shang, "Bridging the divide between software developers and operators using logs," in *ICSE '12: Proceedings of the 34th International Conference on Software Engineering*.
45. W. Shang, Z. M. Jiang, B. Adams, A. E. Hassan, M. W. Godfrey, M. Nasser, and P. Flora, "An exploratory study of the evolution of communicated information about the execution of large software systems," *Journal of Software: Evolution and Process*, vol. 26, no. 1, pp. 3–26, 2014.